

Final Exam Covid Data

D. Clasen

2026-08-17

Overview and Question being asked

This report pulls the daily COVID-19 time series data and builds a few visualizations and a simple linear model looking at the relationship between how widely a state was infected and how deadly the outbreak was there. I wanted to see the states that were the highest case and death rate as well as my own state of Arizona.

Importing the Data

```
url_in <- "https://raw.githubusercontent.com/CSSEGISandData/COVID-19/refs/heads/master/csse_covid_19_data"

file_names <- c("time_series_covid19_confirmed_US.csv",
                "time_series_covid19_confirmed_global.csv",
                "time_series_covid19_deaths_US.csv",
                "time_series_covid19_deaths_global.csv")

urls <- str_c(url_in, file_names)

us_cases      <- read_csv(urls[1])
global_cases  <- read_csv(urls[2])
us_deaths     <- read_csv(urls[3])
global_deaths <- read_csv(urls[4])
```

Tidying the Global Data

Each global file has one row per country and a separate column for every date. The block below pivots those date columns into a single `date` column with a `cases/deaths` value column, drops `Lat/Long` joins the cases and deaths tables together, converts `date` from text to an actual date, and renames the `Province/State` and `Country/Region` columns to underscore-separated names.

```
global_cases <- global_cases %>%
  pivot_longer(cols = -c(`Province/State`, `Country/Region`, Lat, Long),
               names_to = "date", values_to = "cases") %>%
  select(-c(Lat, Long))

global_deaths <- global_deaths %>%
  pivot_longer(cols = -c(`Province/State`, `Country/Region`, Lat, Long),
               names_to = "date", values_to = "deaths") %>%
```

```

select(-c(Lat, Long))

global <- global_cases %>%
  full_join(global_deaths,
            by = c("Province/State", "Country/Region", "date")) %>%
  rename(Country_Region = `Country/Region`,
         Province_State = `Province/State`) %>%
  mutate(date = mdy(date)) %>%
  filter(cases > 0)

```

Adding Population Data to the Global Data

The time series files don't include population, so this brings it in. It also combines the state or province with the country in its own column.

```

global <- global %>%
  unite("Combined_Key", c(Province_State, Country_Region),
        sep = ", ", na.rm = TRUE, remove = FALSE)

uid_lookup_url <- "https://raw.githubusercontent.com/CSSEGISandData/COVID-19/master/csse_covid_19_data/"

uid <- read_csv(uid_lookup_url, show_col_types = FALSE) %>%
  select(-c(Lat, Long_, Combined_Key, code3, iso2, iso3, Admin2))

global <- global %>%
  left_join(uid, by = c("Province_State", "Country_Region")) %>%
  select(-c(UID, FIPS)) %>%
  select(Province_State, Country_Region, date, cases, deaths,
        Population, Combined_Key)

```

Tidying the US Data

The US files had many unneeded rows and columns. This code cleans them up

```

us_cases <- us_cases %>%
  pivot_longer(cols = -(UID:Combined_Key),
               names_to = "date", values_to = "cases") %>%
  select(Admin2:cases) %>%
  mutate(date = mdy(date)) %>%
  select(-c(Lat, Long_))

us_deaths <- us_deaths %>%
  pivot_longer(cols = -(UID:Population),
               names_to = "date", values_to = "deaths") %>%
  select(Admin2:deaths) %>%
  mutate(date = mdy(date)) %>%
  select(-c(Lat, Long_))

us <- us_cases %>%
  full_join(us_deaths,
            by = c("Admin2", "Province_State", "Country_Region",
                  "Combined_Key", "date"))

```

Aggregating to State and National Level

Next one gives one row per state and cleans up the totals


```
us_by_state <- us %>%
  group_by(Province_State, Country_Region, date) %>%
  summarize(cases = sum(cases), deaths = sum(deaths),
            Population = sum(Population), .groups = "drop") %>%
  mutate(deaths_per_mill = deaths * 1000000 / Population) %>%
  select(Province_State, Country_Region, date,
         cases, deaths, deaths_per_mill, Population)

us_totals <- us_by_state %>%
  group_by(Country_Region, date) %>%
  summarize(cases = sum(cases), deaths = sum(deaths),
            Population = sum(Population), .groups = "drop") %>%
  mutate(deaths_per_mill = deaths * 1000000 / Population) %>%
  select(Country_Region, date, cases, deaths, deaths_per_mill, Population)
```

Visualization 1: Cumulative Cases and Deaths in the US

This plot shows the running total of confirmed cases and deaths for the country over time.

```
us_totals %>%
  filter(cases > 0) %>%
  ggplot(aes(x = date, y = cases)) +
  geom_line(aes(color = "cases")) +
  geom_point(aes(color = "cases")) +
  geom_line(aes(y = deaths, color = "deaths")) +
  geom_point(aes(y = deaths, color = "deaths")) +
  scale_y_log10() +
  theme(legend.position = "bottom",
        axis.text.x = element_text(angle = 90)) +
  labs(title = "COVID19 in US", y = NULL)
```




Visualization 2: Cumulative Cases and Deaths for a Single State

Same idea as above, but filtered down to just Arizona

```
state <- "Arizona"

us_by_state %>%
  filter(Province_State == state, cases > 0) %>%
  ggplot(aes(x = date, y = cases)) +
  geom_line(aes(color = "cases")) +
  geom_point(aes(color = "cases")) +
  geom_line(aes(y = deaths, color = "deaths")) +
  geom_point(aes(y = deaths, color = "deaths")) +
  scale_y_log10() +
  theme(legend.position = "bottom",
        axis.text.x = element_text(angle = 90)) +
  labs(title = str_c("COVID19 in ", state), y = NULL)
```

COVID19 in Arizona



Daily New Cases and Deaths

The columns built so far are cumulative totals, which are useful for seeing overall scale but this next code allows to see the change day to day.

```
us_by_state <- us_by_state %>%  
  group_by(Province_State) %>%  
  mutate(new_cases = cases - lag(cases),  
         new_deaths = deaths - lag(deaths)) %>%  
  ungroup()  
  
us_totals <- us_totals %>%  
  mutate(new_cases = cases - lag(cases),  
         new_deaths = deaths - lag(deaths))
```

Visualization 3: Daily New Cases and Deaths

Same log-scale line/point pattern as before, but now showing the daily change rather than the running total. When seeing this you can see that the values begin to decrease as the dates continue forward and historically this makes sense as there was more of a handle on vaccines and quarantine practices. Also you can see just how much more spread out it is in the newer dates compared to the height of Covid in 2020/2021.

```
us_totals %>%
  filter(new_cases > 0) %>%
  ggplot(aes(x = date, y = new_cases)) +
  geom_line(aes(color = "new_cases")) +
  geom_point(aes(color = "new_cases")) +
  geom_line(aes(y = new_deaths, color = "new_deaths")) +
  geom_point(aes(y = new_deaths, color = "new_deaths")) +
  scale_y_log10() +
  theme(legend.position = "bottom",
        axis.text.x = element_text(angle = 90)) +
  labs(title = "COVID19 in US", y = NULL)
```

COVID19 in US



```
us_by_state %>%
  filter(Province_State == state, new_cases > 0) %>%
  ggplot(aes(x = date, y = new_cases)) +
  geom_line(aes(color = "new_cases")) +
  geom_point(aes(color = "new_cases")) +
  geom_line(aes(y = new_deaths, color = "new_deaths")) +
  geom_point(aes(y = new_deaths, color = "new_deaths")) +
  scale_y_log10() +
  theme(legend.position = "bottom",
        axis.text.x = element_text(angle = 90)) +
  labs(title = str_c("COVID19 in ", state), y = NULL)
```

COVID19 in Arizona



State-Level Summary: Rates per Thousand Residents

This collapses each state down to a single row with its maximum values and makes them in rate per thousand.

```
us_state_totals <- us_by_state %>%
  group_by(Province_State) %>%
  summarize(deaths = max(deaths), cases = max(cases),
            population = max(Population),
            cases_per_thou = 1000 * cases / population,
            deaths_per_thou = 1000 * deaths / population) %>%
  filter(cases > 0, population > 0)
```

Exploring the Best and Worst States by Death Rate

A quick look at which states had the lowest and highest deaths per thousand. When looking at this you can see that Arizona had the highest rate per thousand and American Samoa had the least.

```
us_state_totals %>%
  slice_min(deaths_per_thou, n = 10) %>%
  select(deaths_per_thou, cases_per_thou, everything())
```

```
## # A tibble: 10 x 6
```

```
##   deaths_per_thou cases_per_thou Province_State      deaths    cases population
```

```
##           <dbl>           <dbl> <chr>           <dbl> <dbl>           <dbl>
## 1         0.611           150. American Samoa         34 8.32e3         55641
## 2         0.744           248. Northern Mariana Isl~    41 1.37e4         55144
## 3         1.21            231. Virgin Islands         130 2.48e4        107268
## 4         1.30            269. Hawaii           1841 3.81e5       1415872
## 5         1.49            245. Vermont           929 1.53e5        623989
## 6         1.55            293. Puerto Rico        5823 1.10e6       3754939
## 7         1.65            340. Utah             5298 1.09e6       3205958
## 8         2.01            415. Alaska           1486 3.08e5        740995
## 9         2.03            252. District of Columbia  1432 1.78e5        705749
## 10        2.06            253. Washington        15683 1.93e6       7614893
```

```
us_state_totals %>%
  slice_max(deaths_per_thou, n = 10) %>%
  select(deaths_per_thou, cases_per_thou, everything())
```

```
## # A tibble: 10 x 6
##   deaths_per_thou cases_per_thou Province_State deaths   cases population
##           <dbl>           <dbl> <chr>           <dbl>   <dbl>      <dbl>
## 1         4.55           336. Arizona        33102 2443514    7278717
## 2         4.54           326. Oklahoma        17972 1290929    3956971
## 3         4.49           333. Mississippi     13370 990756     2976149
## 4         4.44           359. West Virginia    7960 642760     1792147
## 5         4.32           320. New Mexico       9061 670929     2096829
## 6         4.31           334. Arkansas        13020 1006883    3017804
## 7         4.29           335. Alabama         21032 1644533    4903185
## 8         4.28           368. Tennessee       29263 2515130    6829174
## 9         4.23           307. Michigan        42205 3064125    9986857
## 10        4.06           385. Kentucky        18130 1718471    4467673
```

Linear Model: Deaths per Thousand vs. Cases per Thousand

This model asks a simple question: across states, does a higher case rate predict a higher death rate?

```
mod <- lm(deaths_per_thou ~ cases_per_thou, data = us_state_totals)
summary(mod)
```

```
##
## Call:
## lm(formula = deaths_per_thou ~ cases_per_thou, data = us_state_totals)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3352 -0.5978  0.1491  0.6535  1.2086
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -0.36167    0.72480  -0.499    0.62
## cases_per_thou  0.01133    0.00232   4.881 9.76e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```




```
## Residual standard error: 0.8615 on 54 degrees of freedom
## Multiple R-squared:  0.3061, Adjusted R-squared:  0.2933
## F-statistic: 23.82 on 1 and 54 DF,  p-value: 9.763e-06
```

Visualization 4: Actual vs. Predicted Deaths per Thousand

This plots each state's actual death rate (blue) against the death rate the model predicts from its case rate alone (red). Where the red and blue points for a state are far apart, that state's actual outcome deviated a lot from what its case rate alone would suggest.

```
us_tot_w_pred <- us_state_totals %>%
  mutate(pred = predict(mod))

us_tot_w_pred %>%
  ggplot() +
  geom_point(aes(x = cases_per_thou, y = deaths_per_thou), color = "blue") +
  geom_point(aes(x = cases_per_thou, y = pred), color = "red")
```



Bias

There is bias in the variables we have chosen and how we even cleaned the data. It is important to see that we are deciding all of the variables and what we are actually looking at in each visualization. This was all set up as our own choice to what we wanted to look at. These numbers are also a little outdated because the newest date was in 2023. This means that there are 3 years of missing data.

Conclusion

Overall the information was a good representation of the peak times of covid and gave a good understanding of the states with the highest cases and death rates. I enjoyed looking at Arizona especially because it allowed me to see where I live and the information associated with it.